



Lightning-Fast Cluster Computing

DataFrame and Streaming

13/05/2019 - Big Data 2019

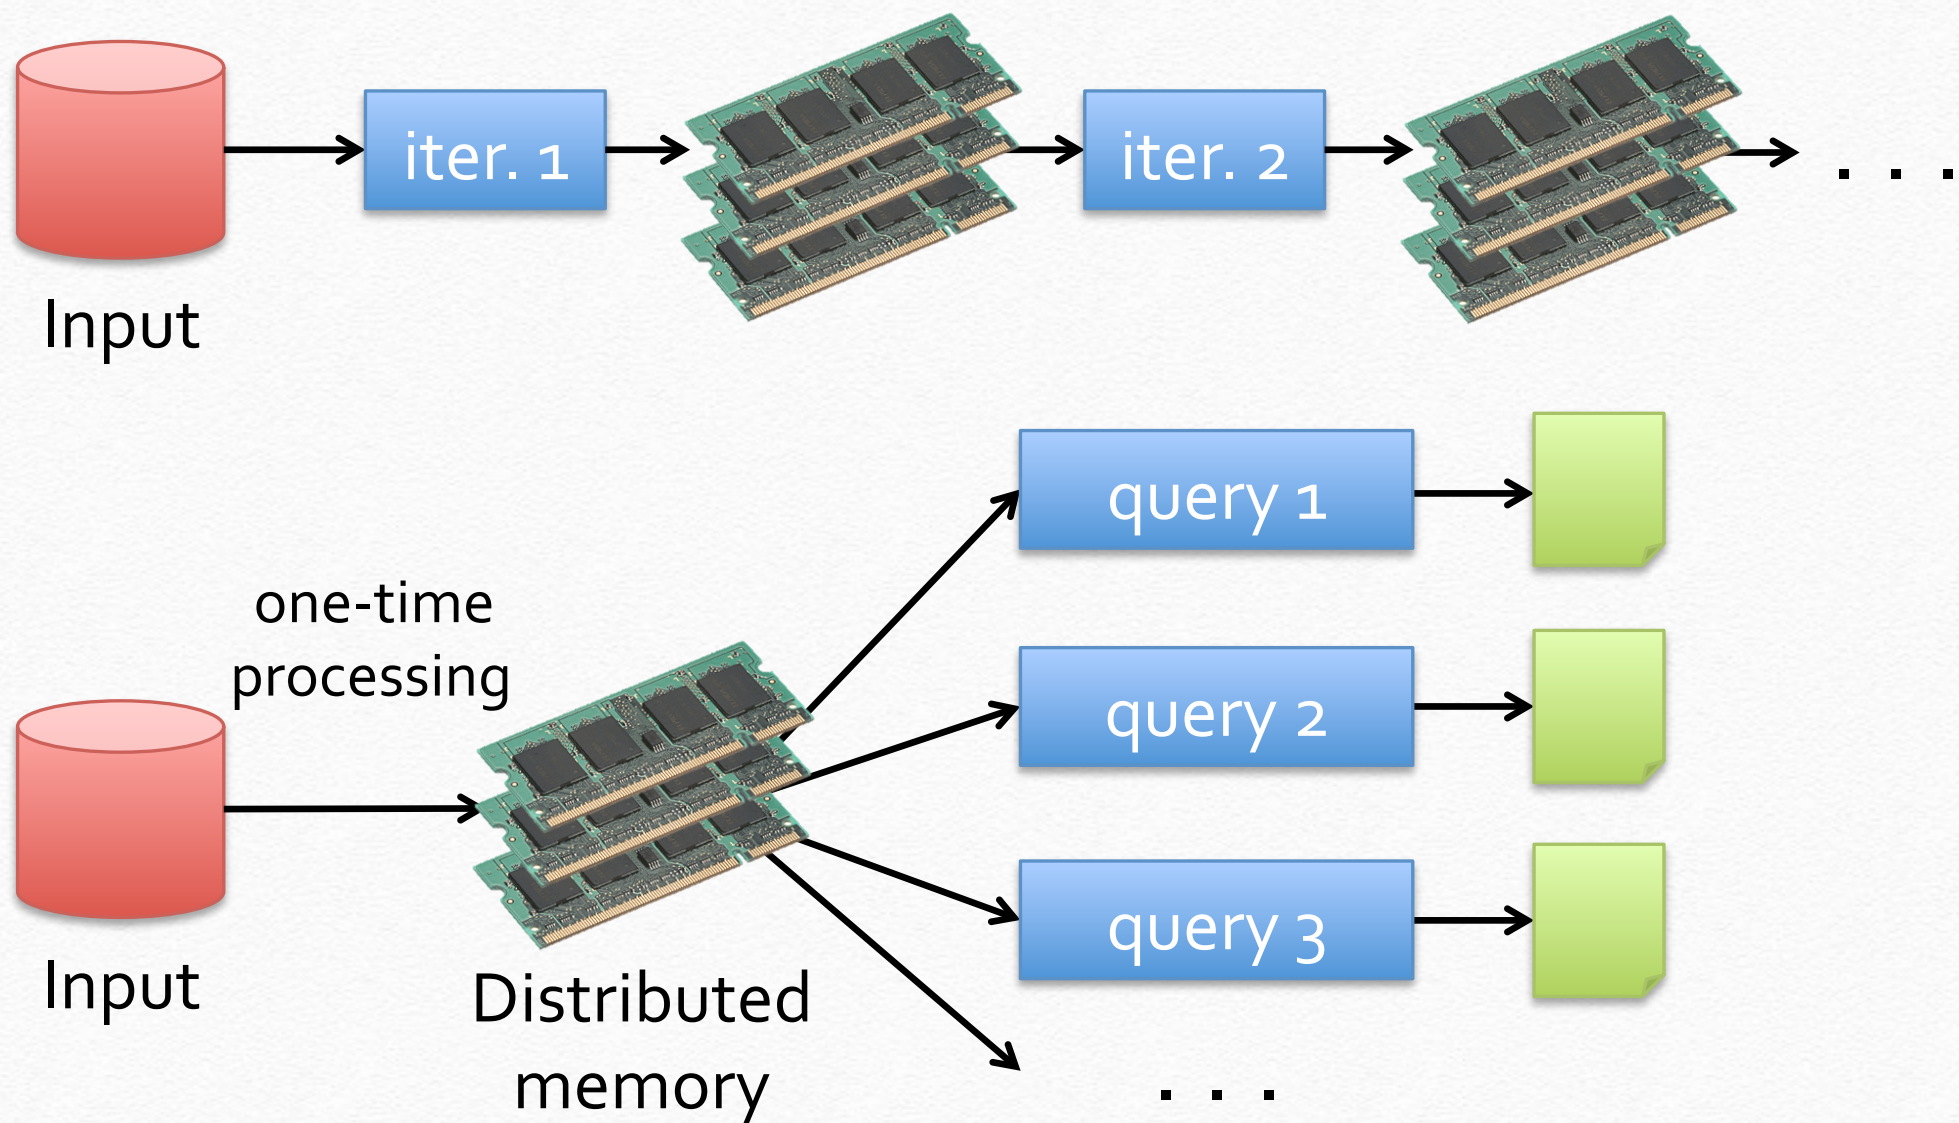


Apache Spark

- ❖ **Not** a modified version of **Hadoop**
- ❖ Separate, fast, **MapReduce-like** engine
 - ☑ **In-memory** data storage for very fast iterative queries
 - ☑ General **execution graphs** and powerful optimizations
 - ☑ Up to **40x faster** than Hadoop
- ❖ **Compatible** with Hadoop's storage APIs
 - ☑ Can read/write to any Hadoop-supported system, including HDFS, HBase, SequenceFiles, etc

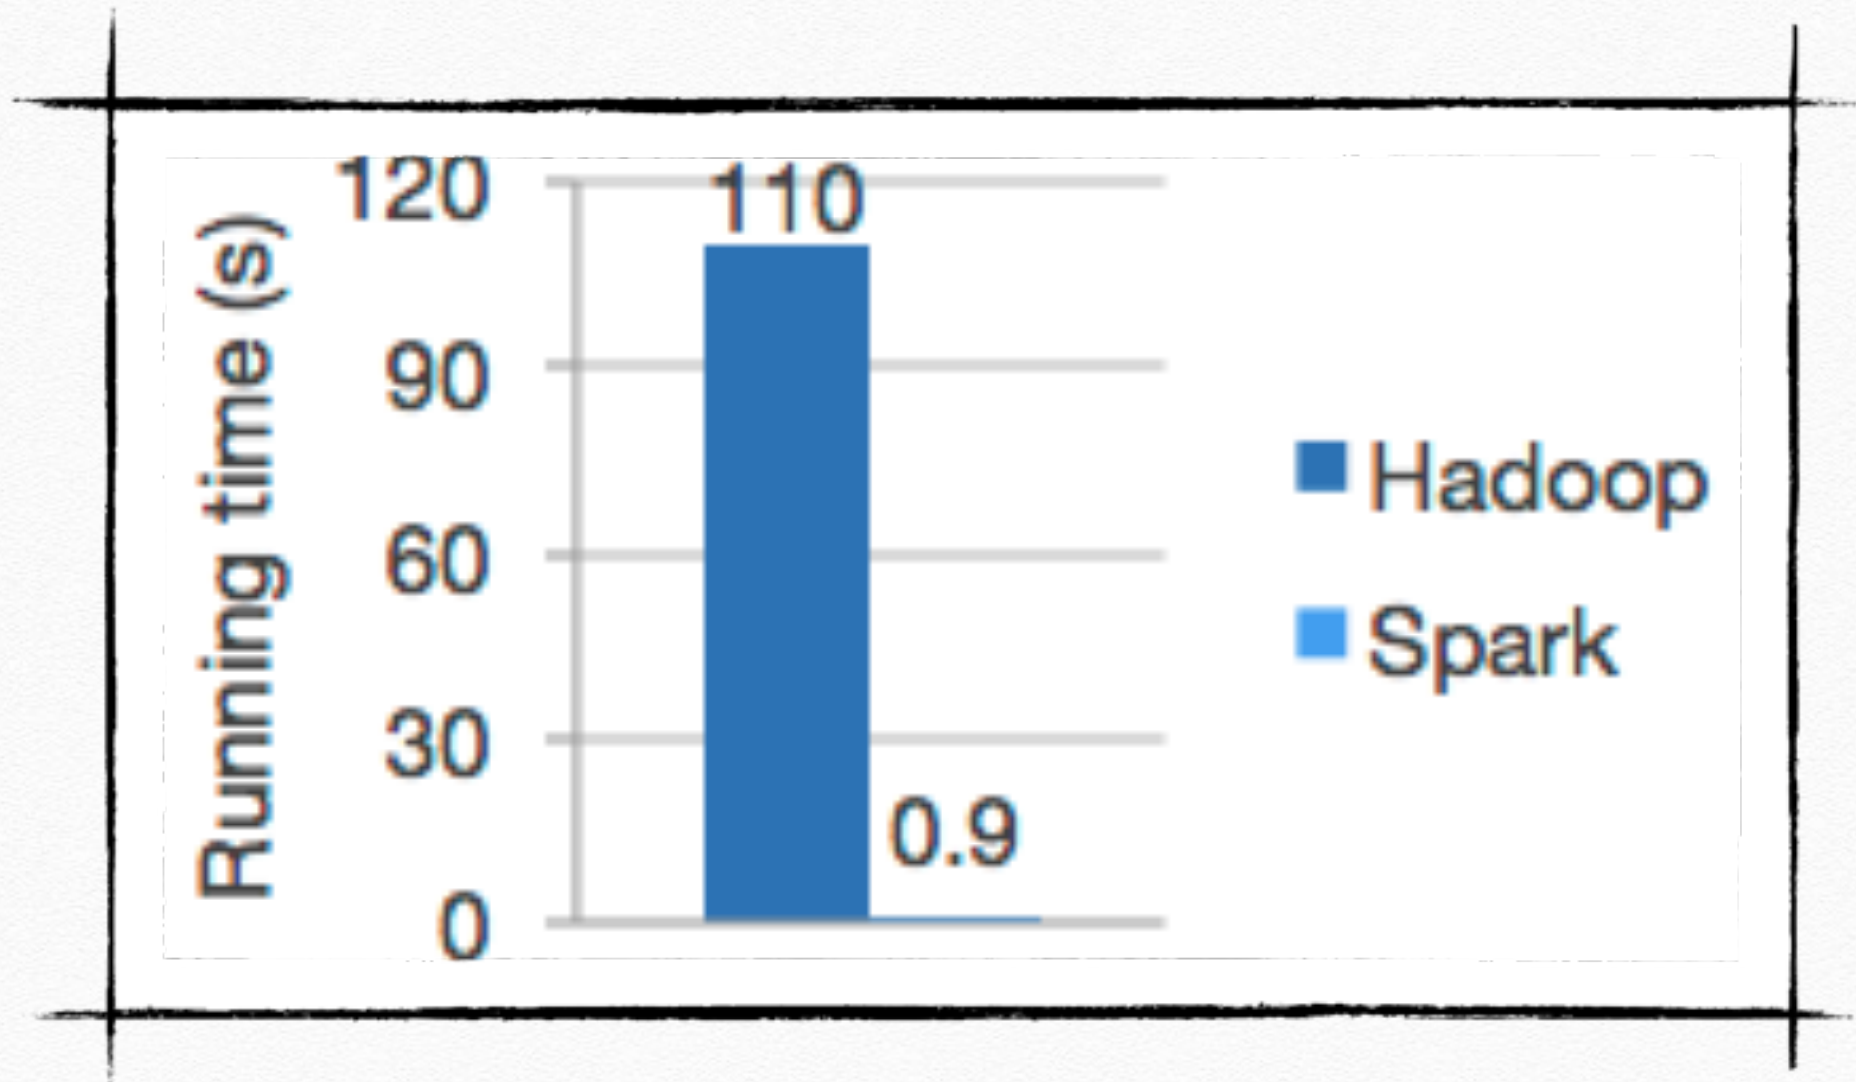


Apache Spark





Apache Spark



10-100× faster than network and disk

Spark

Users

CONVIVA®

foursquare

quantifind

KLOUT

YAHOO!
RESEARCH

University of California
Berkeley

PRINCETON
UNIVERSITY

UCSF



Combine SQL, streaming, and complex analytics

Spark
SQL

Spark
Streaming

MLlib
(machine
learning)

GraphX
(graph)

Apache Spark



Spark Configuration

- ❖ Download a **binary release** of **apache Spark**:
- ❖ **spark-2.3.0-bin-hadoop2.7.tgz**

Download Apache Spark™

1. Choose a Spark release:
2. Choose a package type:
3. Download Spark: [spark-2.3.0-bin-hadoop2.7.tgz](#)
4. Verify this release using the [2.3.0 signatures and checksums](#) and [project release KEYS](#).

Note: Starting version 2.0, Spark is built with Scala 2.11 by default. Scala 2.10 users should download the Spark source package and build [with Scala 2.10 support](#).



Spark Running

- ❖ Running Spark Shell [**scala**]:

```
$:~spark-*/bin/spark-shell
```

- ❖ Running Spark Shell [**python**]:

```
$:~spark-*/bin/pyspark
```

- ❖ **Spark Shell - Scala**

```
Welcome to
```

```
  ____
 /  __/  __  ____  /  /  __
 _\  \/_  \/_  _\  \/_  _\
/_  _/  .__/_\_,/_/_/_/_\_,_  version 1.3.1
  /_/_
```

```
Using Scala version 2.10.4 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_05)
```

```
Type in expressions to have them evaluated.
```

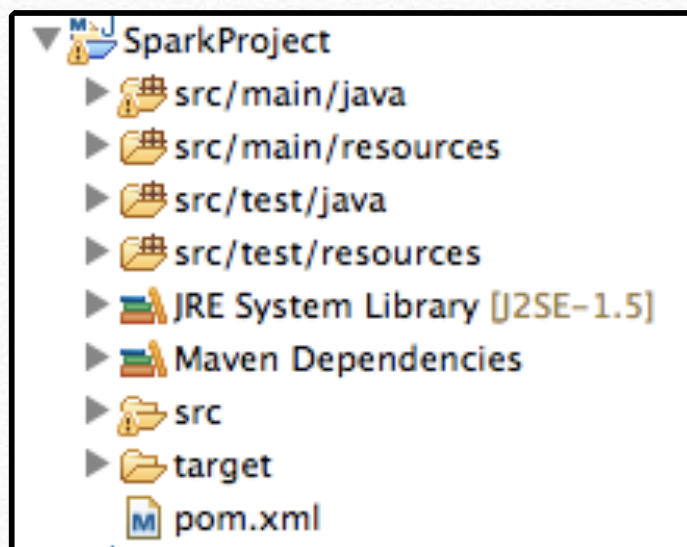
```
scala>
```




Spark Self-contained applications

❖ pom.xml

❖ Maven Project



```
<project>
  <modelVersion>4.0.0</modelVersion>
  <groupId>edu.berkeley</groupId>
  <artifactId>simple-project</artifactId>
  <version>1.0</version>
  <dependencies>
    <dependency> <!-- Spark dependency -->
      <groupId>org.apache.spark</groupId>
      <artifactId>spark-core_2.10</artifactId>
      <version>1.6.1</version>
    </dependency>

    <dependency>
      <groupId>org.apache.spark</groupId>
      <artifactId>spark-streaming_2.10</artifactId>
      <version>1.6.1</version>
    </dependency>
    <dependency>
      <groupId>org.apache.spark</groupId>
      <artifactId>spark-sql_2.10</artifactId>
      <version>1.6.1</version>
    </dependency>
  </dependencies>
</project>
```




Spark Self-contained applications

❖ SimpleApp.java

create logData: an Object like [line1, line2, line3, ...]

sopra la panca la capra campa, sotto la panca la capra crepa

Lines with **a**: 1, lines with **b**: 0



Spark Self-contained applications

❖ Java Spark API

```
import org.apache.spark.api.java.*;

import org.apache.spark.SparkConf;

import org.apache.spark.api.java.function.Function;

public class SimpleApp {

    public static void main(String[] args) {

        String logFile = "data/simpleLog.txt"; // you can use "YOUR_SPARK_HOME/README.md"

        SparkConf conf = new SparkConf().setAppName("Simple Application");

        JavaSparkContext sc = new JavaSparkContext(conf);

        JavaRDD<String> logData = sc.textFile(logFile).cache();

        long numAs = logData.filter(new Function<String, Boolean>() {

            public Boolean call(String s) { return s.contains("a"); }

        }).count();

        long numBs = logData.filter(new Function<String, Boolean>() {

            public Boolean call(String s) { return s.contains("b"); }

        }).count();

        System.out.println("Lines with a: " + numAs + ", lines with b: " + numBs);

    }

}
```




Spark Self-contained applications

❖ Java Spark API: configuration of Spark application

```
String logFile = "data/simpleLog.txt";

SparkConf conf = new SparkConf().setAppName("Simple Application");

JavaSparkContext sc = new JavaSparkContext(conf);

JavaRDD<String> logData = sc.textFile(logFile).cache();
```

Resilient Distributed Datasets (RDD)
is a fundamental data structure of
Spark. It is an immutable distributed
collection of objects.



Spark Self-contained applications

❖ Java Spark API: Spark actions

```
long numAs = logData.filter(new Function<String, Boolean>() {  
    public Boolean call(String s) { return s.contains("a"); }  
}).count();  
  
long numBs = logData.filter(new Function<String, Boolean>() {  
    public Boolean call(String s) { return s.contains("b"); }  
}).count();  
  
System.out.println("Lines with a: " + numAs + ", lines with b: " + numBs);
```




Spark Running - standalone

- ❖ Running Java Spark applications:

```
$:~spark-*/bin/spark-submit
```

```
--class "SimpleApp"
```

```
--master local[4]
```

```
SparkProject-1.0.jar
```

**local to run locally
with one thread, or
local[N] to run locally
with N threads.**

- ❖ output [terminal] - using simpleLog.txt

```
Lines with a: 1, Lines with b: 0
```




Spark Running - standalone

- ❖ Running Java Spark applications:

```
$:~spark-*/bin/spark-submit
```

```
--class "SimpleApp"
```

```
--master local[4]
```

```
SparkProject-1.0.jar
```

**local to run locally
with one thread, or
local[N] to run locally
with N threads.**

- ❖ output [terminal] - using README.md

```
Lines with a: 46, Lines with b: 23
```




Spark Running - YARN

- ❖ Running Java Spark applications:

```
$:~spark-*/bin/spark-submit
```

```
--class "SimpleApp"
```

```
--master yarn
```

```
SparkProject-1.0.jar
```

**The --master option
allows to specify the
master URL for a
distributed cluster**

- ❖ output [terminal] - using README.md

```
Lines with a: 46, Lines with b: 23
```




Spark Running - AWS

❖ Look at

<http://docs.aws.amazon.com/emr/latest/ReleaseGuide/emr-spark-launch.html>



Lambda Expressions

- ❖ You're usually trying to pass functionality as an **argument** to another **method**, such as what action should be taken when someone clicks a button.
- ❖ **Lambda expressions** enable you to do this, to **treat functionality** as **method argument**, or **code as data**



Lambda Expressions

- ❖ Suppose that you are creating a **social networking application**.
- ❖ You want to create a **feature** that enables an administrator to perform **any kind of action**, such as *sending a message, on members of the social networking application that satisfy certain criteria*.
- ❖ Suppose that members of this social networking application are represented by the following **Person** class:

```
public class Person {  
    public enum Sex {  
        MALE, FEMALE  
    }  
  
    String name;  
    LocalDate birthday;  
    Sex gender;  
    String emailAddress;  
  
    public int getAge() {  
        // ...  
    }  
  
    public void printPerson() {  
        // ...  
    }  
}
```




Lambda Expressions

- ❖ Suppose that the members of your social networking application are stored in a `List<Person>` instance
- ❖ **Approach 1:** Create Methods That Search for Members That Match One Characteristic

```
public static void printPersonsOlderThan(List<Person> roster, int age)
{
    for (Person p : roster) {
        if (p.getAge() >= age) {
            p.printPerson();
        }
    }
}
```




Lambda Expressions

- ❖ Suppose that the members of your social networking application are stored in a `List<Person>` instance
- ❖ **Approach 2:** Create More Generalized Search Methods

```
public static void printPersonsWithinAgeRange(List<Person> roster,
                                              int low, int high) {
    for (Person p : roster) {
        if (low <= p.getAge() && p.getAge() < high) {
            p.printPerson();
        }
    }
}
```




Lambda Expressions

❖ Approach 3: Specify Search Criteria Code in a Local Class

```
public static void printPersons(List<Person> roster,
                                CheckPerson tester) {
    for (Person p : roster) {
        if (tester.test(p)) {
            p.printPerson();
        }
    }
}

interface CheckPerson {
    boolean test(Person p);
}

class CheckPersonEligibleForSelectiveService implements CheckPerson {
    public boolean test(Person p) {
        return p.gender == Person.Sex.MALE &&
            p.getAge() >= 18 &&
            p.getAge() <= 25;
    }
}
```




Lambda Expressions

❖ Approach 4: Specify Search Criteria Code in an Anonymous Class

```
printPersons(  
    roster,  
    new CheckPerson() {  
        public boolean test(Person p) {  
            return p.getGender() == Person.Sex.MALE  
                && p.getAge() >= 18  
                && p.getAge() <= 25;  
        }  
    }  
);
```




Lambda Expressions

❖ Approach 5: Specify Search Criteria Code with a Lambda Expression

```
printPersons(  
    roster,  
    (Person p) -> p.getGender() == Person.Sex.MALE  
        && p.getAge() >= 18  
        && p.getAge() <= 25  
);
```



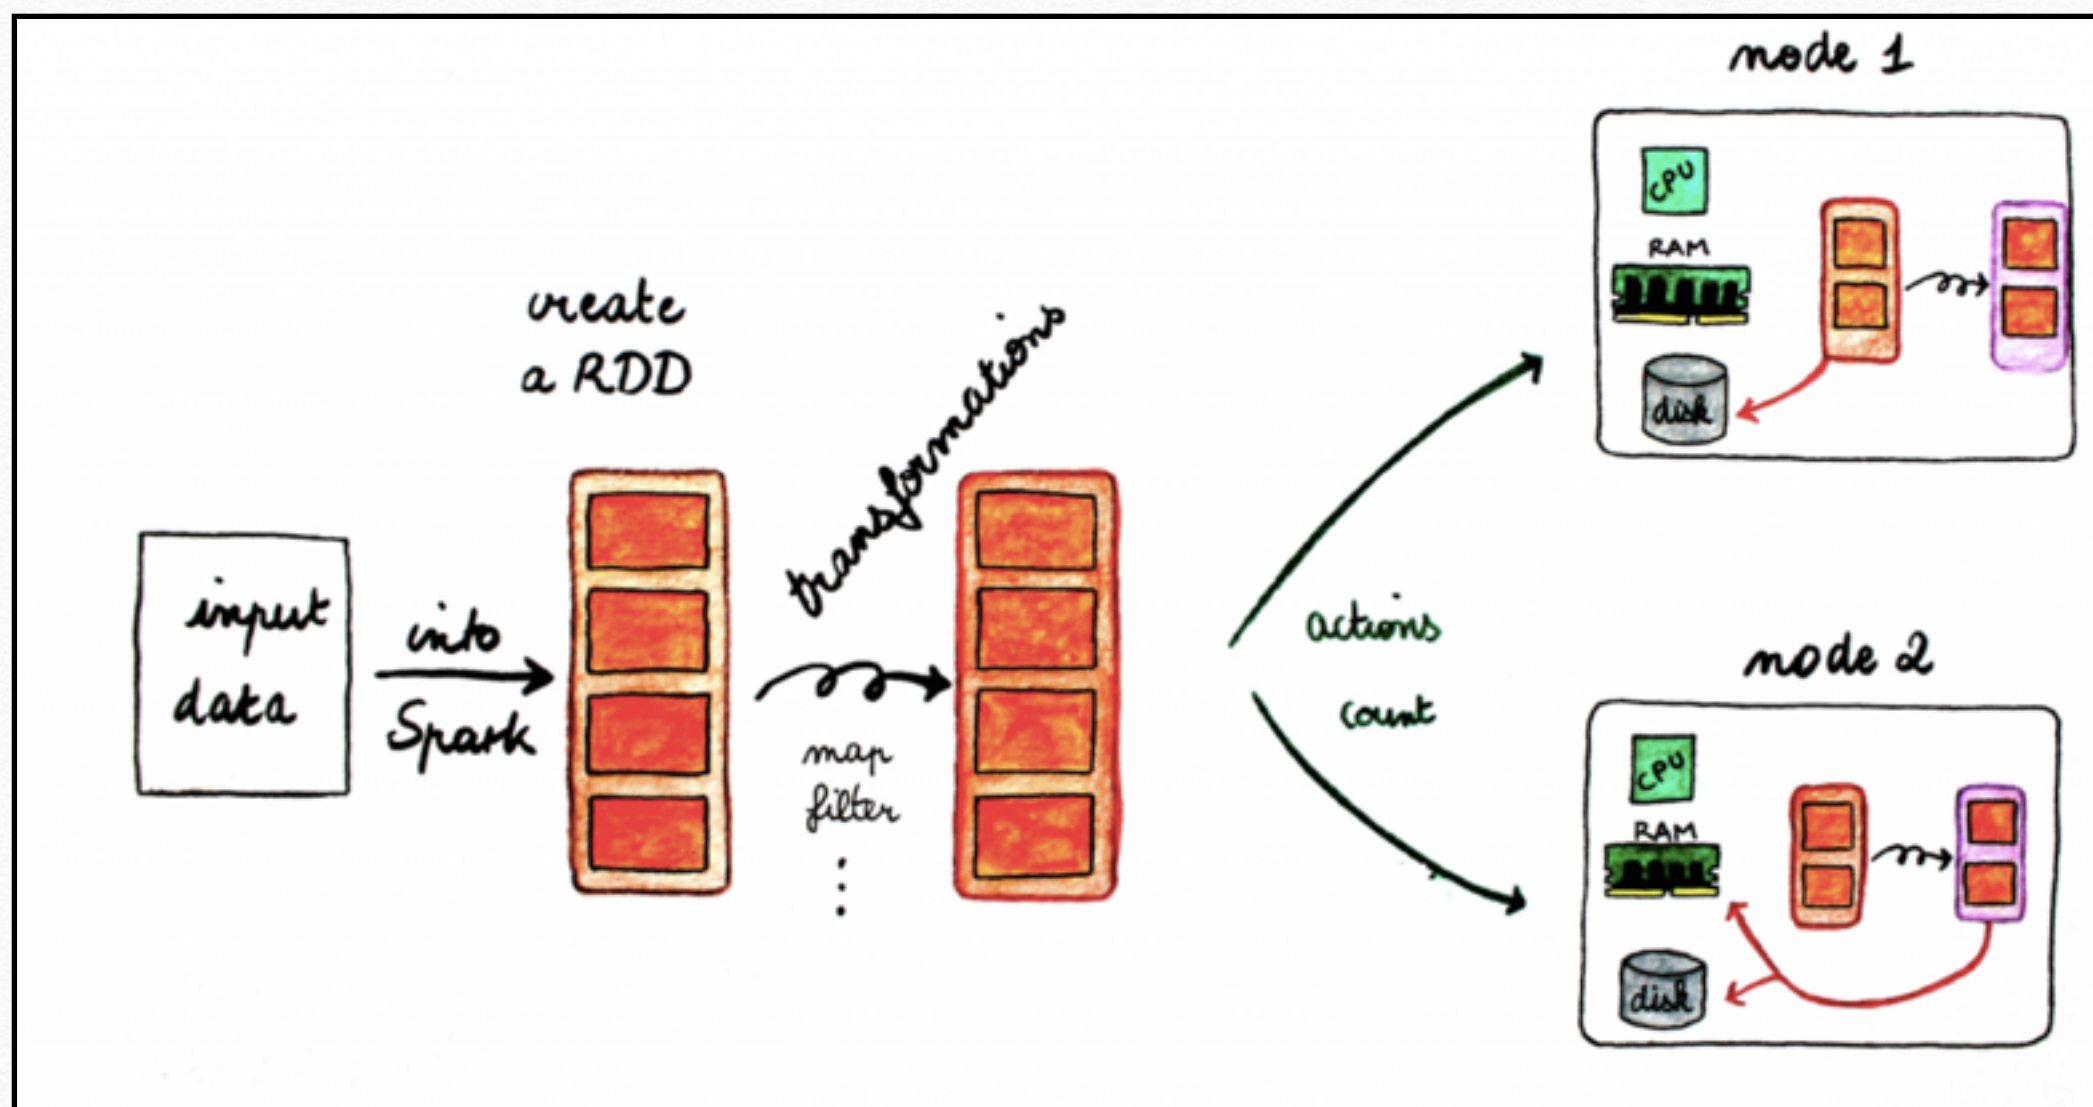

Exercises

❖ http://dia.uniroma3.it/~dvr/es_4.zip



Exercises

❖ SPARK Core API





Exercises

❖ SPARK Core API: Word Count

```
public class Ex0Wordcount {  
  
    private static String pathToFile;  
  
    public Ex0Wordcount(String file){  
        this.pathToFile = file;  
    }  
  
    /**  
     * Load the data from the text file and return an RDD of words  
     */  
    public JavaRDD<String> loadData() {  
  
        SparkConf conf = new SparkConf()  
            .setAppName("Wordcount");  
  
        JavaSparkContext sc = new JavaSparkContext(conf);  
  
        JavaRDD<String> words = sc.textFile(pathToFile)  
            .flatMap(line -> Arrays.asList(line.split(" ")));  
  
        return words;  
  
    }  
}
```




Exercises

❖ SPARK Core API: Word Count

```
public JavaPairRDD<String, Integer> wordcount() {  
    JavaRDD<String> words = loadData();  
  
    // Step 1: mapper step  
    JavaPairRDD<String, Integer> couples =  
        words.mapToPair(word -> new Tuple2<String, Integer>(word, 1));  
  
    // Step 2: reducer step  
    JavaPairRDD<String, Integer> result = couples.reduceByKey((a, b) -> a + b);  
  
    return result;  
}
```




Exercises

❖ SPARK Core API: Word Count

```
/**
 * Now just keep the word which appear strictly more than x times!
 */

public JavaPairRDD<String, Integer> filterOnWordcount(int x) {
    JavaPairRDD<String, Integer> wordcounts = wordcount();

    JavaPairRDD<String, Integer> filtered =
        wordcounts.filter(couple -> couple._2() > x);

    return filtered;
}
```




Exercises

❖ SPARK Core API: Word Count

```
public static void main(String[] args) {  
    // TODO Auto-generated method stub  
  
    if (args.length < 1) {  
        System.err.println("Usage: Ex0Wordcount <filetxt>");  
        System.exit(1);  
    }  
  
    Ex0Wordcount wc = new Ex0Wordcount(args[0]);  
  
    System.out.println("wordcount: "+wc.wordcount().toString());  
}
```




Exercises

❖ SPARK Core API: Tweet Mining

Now we use a dataset with 8198 tweets.
Here an example of a tweet (json):

```
{"id": "572692378957430785",  
  "user": "Sr kian_nishu :)",  
  "text": "@always_nidhi @YouTube no i dnt understand bt i  
          loved of this mve is rocking",  
  "place": "Orissa",  
  "country": "India"  
}
```

We want to make some computations on the tweets:

- Find all the persons mentioned on tweets
- Count how many times each person is mentioned
- Find the 10 most mentioned persons by descending order



Exercises

❖ SPARK Core API: Tweet Mining

```
public class TweetMining {  
  
    private String pathToFile;  
  
    public TweetMining(String file){  
        this.pathToFile = file;  
    }  
  
    // Load the data from the text file and return an RDD of Tweet  
  
    public JavaRDD<Tweet> loadData() { }  
  
    // Find all the persons mentioned on tweets  
  
    public JavaRDD<String> mentionOnTweet() { }  
  
    // Count how many times each person is mentioned  
  
    public JavaPairRDD<String, Integer> countMentions() { }  
  
    // Find the 10 most mentioned persons by descending order  
  
    public List<Tuple2<Integer, String>> top10mentions() { }  
}
```




Exercises

❖ SPARK Core API: Tweet Mining

```
public class Tweet implements Serializable {  
  
    long id; String user; String userName; String text;  
    String place; String country; String lang;  
  
    public String getUserName() { return userName; }  
  
    public String getLang() { return lang; }  
  
    public long getId() { return id; }  
  
    public String getUser() { return user;}  
  
    public String getText() { return text; }  
  
    public String getPlace() { return place; }  
  
    public String getCountry() { return country; }  
  
    @Override  
    public String toString(){  
        return getId() + ", " + getUser() + ", " + getText() + ", " + getPlace() + ", " +  
            getCountry();  
    }  
}
```




Exercises

❖ SPARK Core API: Tweet Mining

```
import com.fasterxml.jackson.databind.ObjectMapper;

public class Parse {

    public static Tweet parseJsonToTweet(String jsonLine) {

        ObjectMapper objectMapper = new ObjectMapper();
        Tweet tweet = null;

        try {
            tweet = objectMapper.readValue(jsonLine, Tweet.class);
        } catch (IOException e) {
            e.printStackTrace();
        }
        return tweet;
    }
}
```




Exercises

❖ SPARK Core API: Tweet Mining (Java 1.7 or later)

```
public JavaRDD<Tweet> loadData() {  
    // create spark configuration and spark context  
    SparkConf conf = new SparkConf()  
        .setAppName("Tweet mining");  
    // .setMaster("local[*]");  
  
    JavaSparkContext sc = new JavaSparkContext(conf);  
  
    JavaRDD<Tweet> tweets = sc.textFile(pathToFile)  
        .map(new Function<String, Tweet>() {  
  
            @Override  
            public Tweet call(String line) throws Exception  
            {  
                return Parse.parseJsonToTweet(line);  
            }  
        });  
  
    return tweets;  
}
```




Exercises

❖ SPARK Core API: Tweet Mining (LAMBDA Java 1.8)

```
public JavaRDD<Tweet> loadData() {  
    // create spark configuration and spark context  
    SparkConf conf = new SparkConf()  
        .setAppName("Tweet mining");  
    // .setMaster("local[*]");  
  
    JavaSparkContext sc = new JavaSparkContext(conf);  
  
    JavaRDD<Tweet> tweets = sc.textFile(pathToFile)  
        .map(line -> Parse.parseJsonToTweet(line));  
  
    return tweets;  
}
```




Exercises

❖ SPARK Core API: Tweet Mining (Java 1.7 or later)

- Find all the persons mentioned on tweets

```
public JavaRDD<String> mentionOnTweet() {
    JavaRDD<Tweet> tweets = loadData();

    JavaRDD<String> mentions = tweets.flatMap(new FlatMapFunction<Tweet,
String>() {
        @Override
        public Iterable<String> call(Tweet tweet) throws Exception {
            return Arrays.asList(tweet.getText().split(" "));
        }
    })
    .filter(new Function<String, Boolean>() {
        @Override
        public Boolean call(String word) throws Exception {
            return word.startsWith("@") && word.length() > 1;
        }
    });

    System.out.println("mentions.count() " + mentions.count());
    return mentions;
}
```




Exercises

❖ SPARK Core API: Tweet Mining (Java 1.8)

- Find all the persons mentioned on tweets

```
public JavaRDD<String> mentionOnTweet() {  
    JavaRDD<Tweet> tweets = loadData();  
  
    JavaRDD<String> mentions =  
        tweets.flatMap(tweet -> Arrays.asList(tweet.getText()  
            .split(" ")))  
            .filter(word -> word.startsWith("@") && word.length() > 1);  
  
    System.out.println("mentions.count() " + mentions.count());  
    return mentions;  
}
```




Exercises

❖ SPARK Core API: Tweet Mining (Java 1.7 or later)

- Count how many times each person is mentioned

```
public JavaPairRDD<String, Integer> countMentions() {
    JavaRDD<String> mentions = mentionOnTweet();

    JavaPairRDD<String, Integer> mentionCount = mentions.mapToPair(new
PairFunction<String, String, Integer>() {
        @Override
        public Tuple2<String, Integer> call(String mention) throws Exception {
            return new Tuple2<>(mention, 1);
        }
    })
    .reduceByKey(new Function2<Integer, Integer, Integer>() {
        @Override
        public Integer call(Integer x, Integer y) throws Exception {
            return x + y;
        }
    });

    return mentionCount;
}
```




Exercises

❖ SPARK Core API: Tweet Mining (Java 1.8)

- Count how many times each person is mentioned

```
public JavaPairRDD<String, Integer> countMentions() {  
    JavaRDD<String> mentions = mentionOnTweet();  
  
    JavaPairRDD<String, Integer> mentionCount =  
        mentions.mapToPair(mention -> new Tuple2<>(mention, 1))  
                .reduceByKey((x, y) -> x + y);  
    return mentionCount;  
}
```




Exercises

❖ SPARK Core API: Tweet Mining (Java 1.7 or later)

- Find the 10 most mentioned persons by descending order

```
public List<Tuple2<Integer, String>> top10mentions() {
    JavaPairRDD<String, Integer> counts = countMentions();

    List<Tuple2<Integer, String>> mostMentioned =
        counts.mapToPair(new PairFunction<Tuple2<String, Integer>, Integer,
String>() {
            @Override
            public Tuple2<Integer, String> call(Tuple2<String, Integer> pair) throws
Exception {
                return new Tuple2<>(pair._2(), pair._1());
            }
        })

        .sortByKey(false)
        .take(10);

    return mostMentioned;
}
```




Exercises

❖ SPARK Core API: Tweet Mining (Java 1.8)

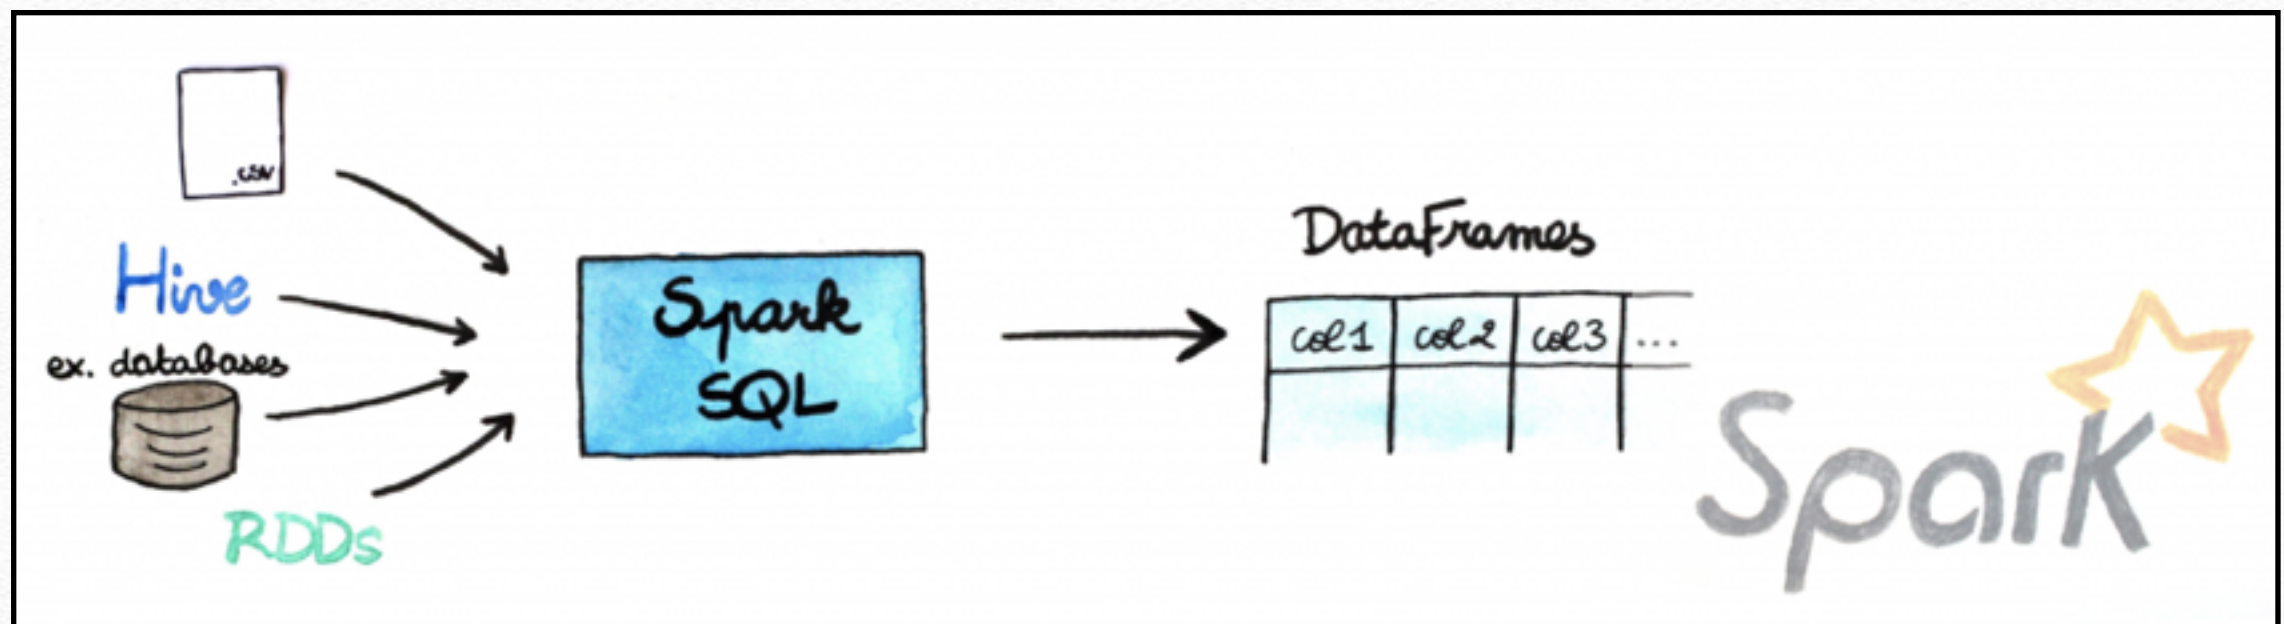
- Find the 10 most mentioned persons by descending order

```
public List<Tuple2<Integer, String>> top10mentions() {  
    JavaPairRDD<String, Integer> counts = countMentions();  
  
    List<Tuple2<Integer, String>> mostMentioned =  
        counts.mapToPair(pair -> new Tuple2<>(pair._2(), pair._1()))  
                .sortByKey(false)  
                .take(10);  
  
    return mostMentioned;  
}
```




Exercises

❖ SPARK SQL (DataFrame)





Exercises

Michael, 29
Andy, 30
Justin, 19

❖ SPARK DataFrame: SparkSQL

```
public class SparkSQL {  
    public static class Person implements Serializable {  
        private String name;  
        private int age;  
  
        public String getName() {  
            return name;  
        }  
  
        public void setName(String name) {  
            this.name = name;  
        }  
  
        public int getAge() {  
            return age;  
        }  
  
        public void setAge(int age) {  
            this.age = age;  
        }  
    }  
}
```




Exercises

❖ SPARK DataFrame: SparkSQL

```
public static void main(String[] args) throws Exception {  
    if (args.length < 2) {  
        System.err.println("Usage: JavaSparkSQL <filetxt> <filejson>");  
        System.exit(1);  
    }  
  
    SparkConf sparkConf = new SparkConf().setAppName("JavaSparkSQL");  
    JavaSparkContext ctx = new JavaSparkContext(sparkConf);  
    SQLContext sqlContext = new SQLContext(ctx);  
}
```

Michael, 29
Andy, 30
Justin, 19

```
{"name":"Michael"}  
{"name":"Andy", "age":30}  
{"name":"Justin", "age":19}
```



Exercises

Michael, 29 Andy, 30 Justin, 19

❖ SPARK DataFrame: SparkSQL

```
System.out.println("=== Data source: RDD ===");
// Load a text file and convert each line to a Java Bean.
JavaRDD<Person> people = ctx.textFile(args[0]).map(
    new Function<String, Person>() {
        @Override
        public Person call(String line) {
            String[] parts = line.split(",");

            Person person = new Person();
            person.setName(parts[0]);
            person.setAge(Integer.parseInt(parts[1].trim()));

            return person;
        }
    });
```




Exercises

Michael, 29 Andy, 30 Justin, 19

❖ SPARK DataFrame: SparkSQL

```
// Apply a schema to an RDD of Java Beans and register it as a table.
DataFrame schemaPeople = sqlContext.createDataFrame(people, Person.class);
schemaPeople.registerTempTable("people");

// SQL can be run over RDDs that have been registered as tables.
DataFrame teenagers =
    sqlContext.sql("SELECT name FROM people WHERE age >= 13 AND age <= 19");

// The results of SQL queries are DataFrames and support all the normal RDD operations.
// The columns of a row in the result can be accessed by ordinal.
List<String> teenagerNames = teenagers.toJavaRDD().map(new Function<Row, String>() {
    @Override
    public String call(Row row) {
        return "Name: " + row.getString(0);
    }
}).collect();
for (String name: teenagerNames) {
    System.out.println(name);
}
```



Exercises

❖ SPARK DataFrame: SparkSQL

```
{"name":"Michael"}  
{"name":"Andy", "age":30}  
{"name":"Justin", "age":19}
```

```
System.out.println("=== Data source: JSON Dataset ===");  
    // A JSON dataset is pointed by path.  
    // The path can be either a single text file or a directory storing text  
files.  
    String path = args[1];  
    // Create a DataFrame from the file(s) pointed by path  
    DataFrame peopleFromJsonFile = sqlContext.read().json(path);  
  
    // Because the schema of a JSON dataset is automatically inferred, to  
write queries,  
    // it is better to take a look at what is the schema.  
    peopleFromJsonFile.printSchema();  
    // The schema of people is ...  
    // root  
    // |-- age: IntegerType  
    // |-- name: StringType
```




Exercises

❖ SPARK DataFrame: SparkSQL

```
{"name":"Michael"}  
{"name":"Andy", "age":30}  
{"name":"Justin", "age":19}
```

```
// Register this DataFrame as a table.  
peopleFromJsonFile.registerTempTable("people");  
  
// SQL statements can be run by using the sql methods provided by sqlContext.  
DataFrame teenagers3 =  
    sqlContext.sql("SELECT name FROM people WHERE age >= 13 AND age <= 19");  
  
// The results of SQL queries are DataFrame and support all the normal RDD  
operations.  
// The columns of a row in the result can be accessed by ordinal.  
teenagerNames = teenagers3.toJavaRDD().map(new Function<Row, String>() {  
    @Override  
    public String call(Row row) { return "Name: " + row.getString(0); }  
}).collect();  
for (String name: teenagerNames) {  
    System.out.println(name);  
}
```



Exercises

❖ SPARK DataFrame: SparkSQL

```
{"name":"Michael"}  
{"name":"Andy", "age":30}  
{"name":"Justin", "age":19}
```

```
// Alternatively, a DataFrame can be created for a JSON dataset represented by  
// a RDD[String] storing one JSON object per string.  
List<String> jsonData = Arrays.asList(  
  "{ \"name\": \"Yin\", \"address\": { \"city\": \"Columbus\", \"state\": \"Ohio\" } }\" );  
JavaRDD<String> anotherPeopleRDD = ctx.parallelize(jsonData);  
DataFrame peopleFromJsonRDD = sqlContext.read().json(anotherPeopleRDD.rdd());  
  
// Take a look at the schema of this new DataFrame.  
peopleFromJsonRDD.printSchema();  
// The schema of anotherPeople is ...  
// root  
// |-- address: StructType  
// |   |-- city: StringType  
// |   |-- state: StringType  
// |-- name: StringType
```




Exercises

❖ SPARK DataFrame: SparkSQL

```
{"name":"Michael"}  
{"name":"Andy", "age":30}  
{"name":"Justin", "age":19}
```

```
peopleFromJsonRDD.registerTempTable("people2");  
  
DataFrame peopleWithCity = sqlContext.sql("SELECT name, address.city FROM people2");  
List<String> nameAndCity = peopleWithCity.toJavaRDD().map(new Function<Row, String>() {  
    @Override  
    public String call(Row row) {  
        return "Name: " + row.getString(0) + ", City: " + row.getString(1);  
    }  
}).collect();  
for (String name: nameAndCity) {  
    System.out.println(name);  
}  
  
ctx.stop();  
}
```



Exercises

❖ SPARK DataFrame: DataFrameOnTweets

We use again a dataset with 8198 tweets.
Here an example of a tweet (json):

```
{"id": "572692378957430785",  
  "user": "Sr kian_nishu :)",  
  "text": "@always_nidhi @YouTube no i dnt understand bt i  
          loved of this mve is rocking",  
  "place": "Orissa",  
  "country": "India"  
}
```

In the exercise we will create a dataframe with the content of a JSON file. We want to:

- print the dataframe
- print the schema of the dataframe
- find people who are located in Paris
- find the user who tweets the more



Exercises

❖ SPARK DataFrame: DataFrameOnTweets

```
public class DataFrameOnTweets {  
  
    //private static String pathToFile = "data/reduced-tweets.json";  
  
    private static String pathToFile;  
  
    public DataFrameOnTweets(String file){  
        this.pathToFile = file;  
    }  
  
    public DataFrame loadData() {  
        SparkConf conf = new SparkConf()  
            .setAppName("Dataframe");  
  
        JavaSparkContext sc = new JavaSparkContext(conf);  
        // Create a sql context: the SQLContext wraps the SparkContext, and is specific to Spark SQL.  
        // It is the entry point in Spark SQL.  
        SQLContext sqlContext = new SQLContext(sc);  
  
        DataFrame dataframe = sqlContext.read().json(pathToFile);  
  
        return dataframe;  
    }  
}
```



Exercises

❖ SPARK DataFrame: DataFrameOnTweets

```
/**
 * See how looks the dataframe
 */
public void showDataFrame() {
    DataFrame dataframe = loadData();
    // Displays the content of the DataFrame to stdout
    dataframe.show();
}

/**
 * Print the schema
 */
public void printSchema() {
    DataFrame dataframe = loadData();

    dataframe.printSchema();
}
```




Exercises

❖ SPARK DataFrame: DataFrameOnTweets

```
/**
 * Find people who are located in Paris
 */
public DataFrame filterByLocation() {
    DataFrame dataframe = loadData();

    DataFrame filtered = dataframe.filter(dataframe.col("place").equalTo("Paris")).toDF();

    return filtered;
}

/**
 * Find the user who tweets the more
 */
public Row mostPopularTwitterer() {
    DataFrame dataframe = loadData();


    // group the tweets by user first
    DataFrame countByUser = dataframe.groupBy(dataframe.col("user")).count();
    // sort by descending order and take the first one
    JavaRDD<Row> result = countByUser.javaRDD().sortBy(x -> x.get(1), false, 1);

    return result.first();
}
```



Exercises

❖ SPARK Streaming





Exercises

❖ SPARK Streaming: SparkSQLStreaming

```
/**
 * Use DataFrames and SQL to count words in UTF8 encoded, '\n'
 delimited text received from the
 * network every second.
 *
 * Usage: JavaSqlNetworkWordCount <hostname> <port>
 * <hostname> and <port> describe the TCP server that Spark
 Streaming would connect to receive data.
 *
 * To run this on your local machine, you need to first run a
 Netcat server
 *      `$ nc -lk 9999`
 * and then run the example
 *      `$ SparkSQLStreaming localhost 9999`
 */
```



Exercises

❖ SPARK Streaming: SparkSQLStreaming

```
/** Java Bean class to be used with the example JavaSqlNetworkWordCount. */  
public class JavaRecord implements java.io.Serializable {  
    private String word;  
  
    public String getWord() {  
        return word;  
    }  
  
    public void setWord(String word) {  
        this.word = word;  
    }  
}
```




Exercises

❖ SPARK Streaming: SparkSQLStreaming

```
public final class SparkSQLStreaming {  
    private static final Pattern SPACE = Pattern.compile(" ");  
  
    public static void main(String[] args) {  
        if (args.length < 2) {  
            System.err.println("Usage: SparkSQLStreaming <hostname> <port>");  
            System.exit(1);  
        }  
  
        //StreamingExamples.setStreamingLogLevels();  
  
        // Create the context with a 1 second batch size  
        SparkConf sparkConf = new SparkConf().setAppName("SparkSQLStreaming");  
        JavaStreamingContext ssc =  
            new JavaStreamingContext(sparkConf, Durations.seconds(1));  
    }  
}
```



Exercises

❖ SPARK Streaming: SparkSQLStreaming

```
// Create a JavaReceiverInputDStream on target ip:port and count the
// words in input stream of \n delimited text (eg. generated by 'nc')
// Note that no duplication in storage level only for running locally.
// Replication necessary in distributed scenario for fault tolerance.
JavaReceiverInputDStream<String> lines =
    ssc.socketTextStream(args[0],
                        Integer.parseInt(args[1]),
                        StorageLevels.MEMORY_AND_DISK_SER);

JavaDStream<String> words =
    lines.flatMap(new FlatMapFunction<String, String>() {
        @Override
        public Iterable<String> call(String x) {
            return Arrays.asList(SPACE.split(x));
        }
    });
```




Exercises

❖ SPARK Streaming: SparkSQLStreaming

```
// Convert RDDs of the words DStream to DataFrame and run SQL query
words.foreachRDD(new VoidFunction2<JavaRDD<String>, Time>() {
    @Override
    public void call(JavaRDD<String> rdd, Time time) {
        SQLContext sqlContext = JavaSQLContextSingleton.getInstance(rdd.context());

        // Convert JavaRDD[String] to JavaRDD[bean class] to DataFrame
        JavaRDD<JavaRecord> rowRDD = rdd.map(new Function<String, JavaRecord>() {
            @Override
            public JavaRecord call(String word) {
                JavaRecord record = new JavaRecord();
                record.setWord(word);
                return record;
            }
        });
    }
});
```



Exercises

❖ SPARK Streaming: SparkSQLStreaming

```
DataFrame wordsDataFrame = sqlContext.createDataFrame(rowRDD, JavaRecord.class);

// Register as table
wordsDataFrame.registerTempTable("words");

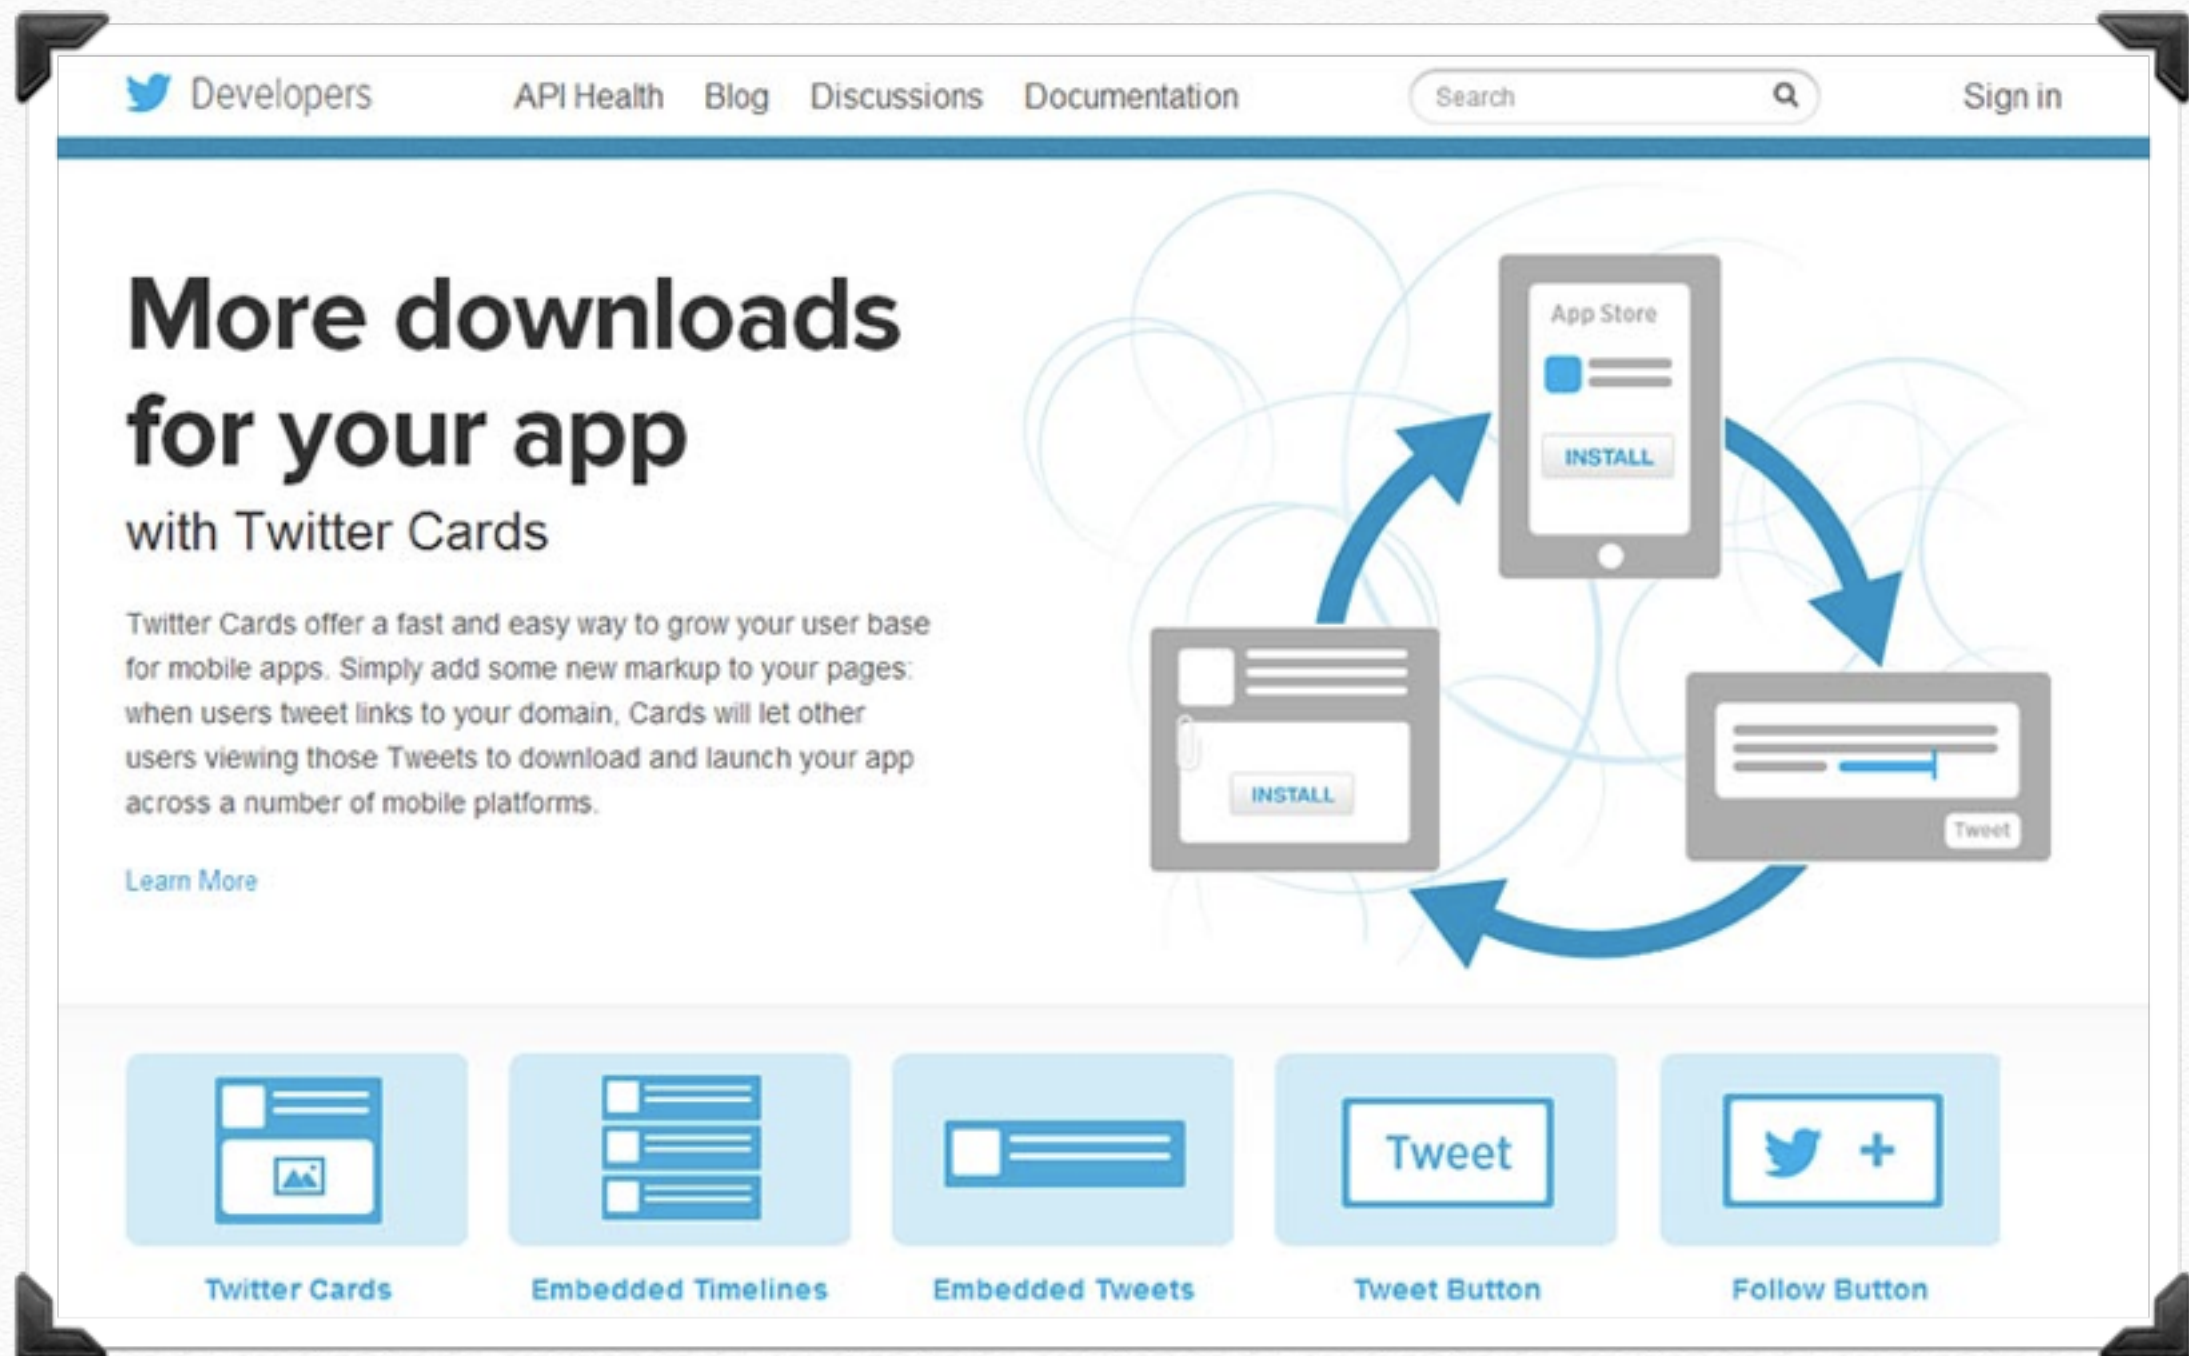
// Do word count on table using SQL and print it
DataFrame wordCountsDataFrame =
    sqlContext.sql("select word, count(*) as total from words group by word");
System.out.println("===== " + time + "=====");
wordCountsDataFrame.show();
    }
});

ssc.start();
ssc.awaitTermination();
}
}
```




How to Register a Twitter App in 8 Easy Steps

❖ #1 Visit the Twitter Developers' Site





How to Register a Twitter App in 8 Easy Steps

❖ #2 Sign in with your Twitter Account

The screenshot shows the Twitter Developers sign-in page. At the top, there is a navigation bar with links for Developers, API Health, Blog, Discussions, and Documentation. A search bar and a 'Sign in' link are also present. Below the navigation bar, the page title 'Sign in with your Twitter account' is displayed. The sign-in form includes a 'Username:' field with the text 'iagdotme' and a 'Password:' field with masked characters. A 'Log in' button is located at the bottom of the form. A link for 'New to Twitter? Sign up' is also visible.

Developers API Health Blog Discussions Documentation Search Sign in

Home

Sign in with your Twitter account

Username: *

iagdotme *

New to Twitter? [Sign up](#)

Password: *

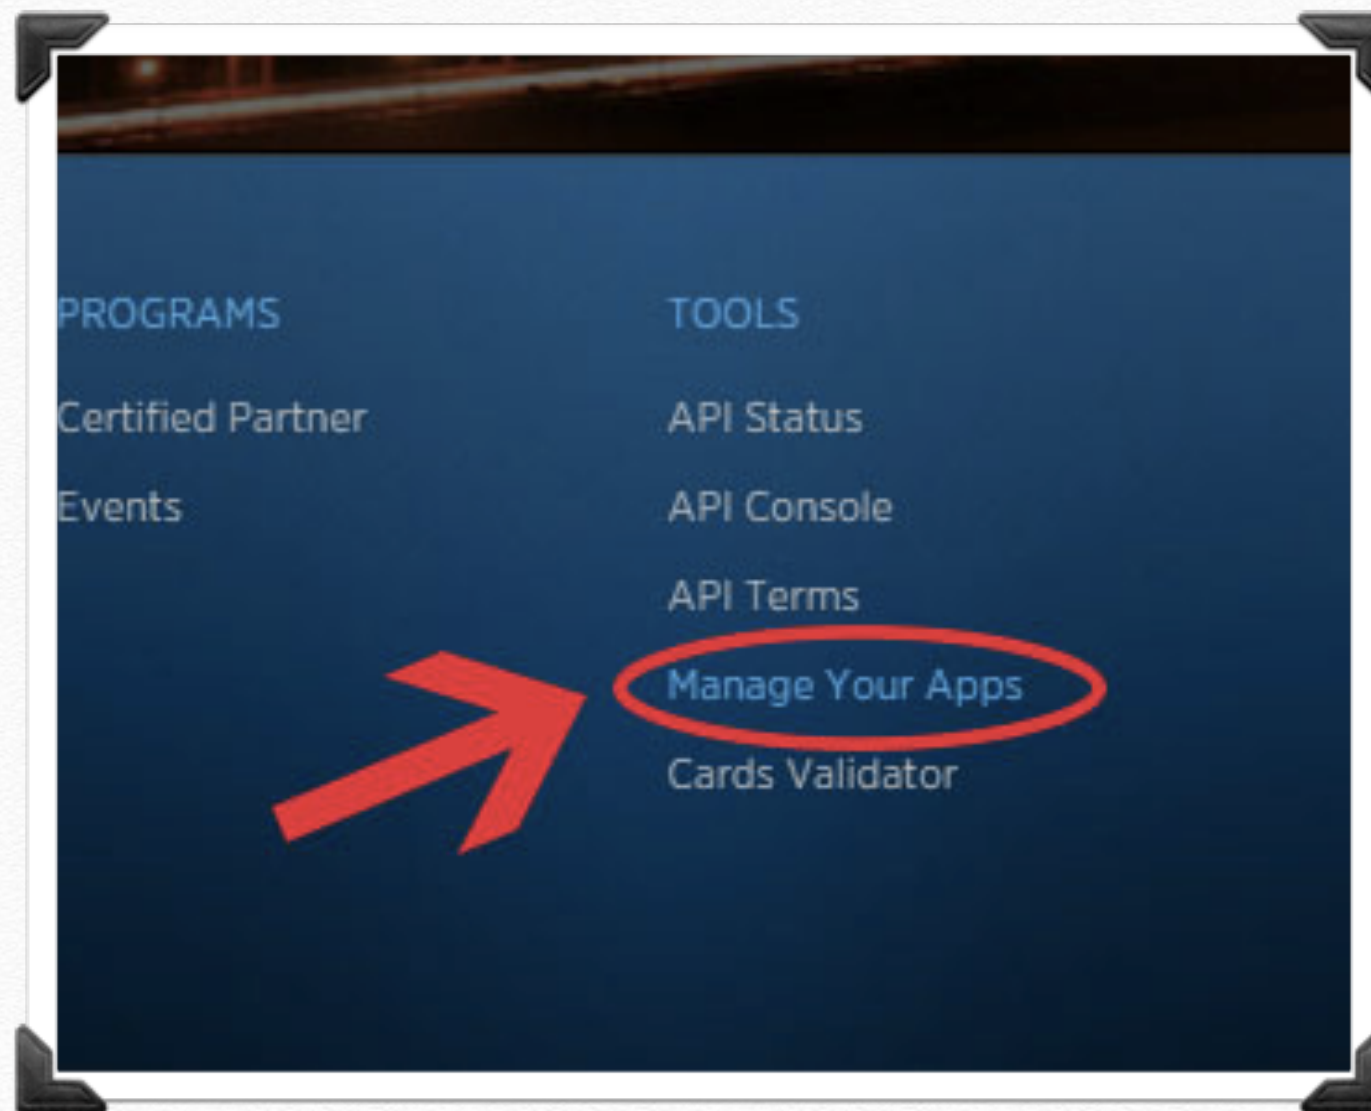
..... *

Log in



How to Register a Twitter App in 8 Easy Steps


❖ #3 Go to apps.twitter.com





How to Register a Twitter App in 8 Easy Steps

❖ #4 Create a New Application





How to Register a Twitter App in 8 Easy Steps

❖ #5 Fill in your Application Details

Home → My applications

Create an application

Application Details

Name: *

Your application name. This is used to attribute the source of a tweet and in user-facing authorization screens. 32 characters max.

Description: *

Your application description, which will be shown in user-facing authorization screens. Between 10 and 200 characters max.

Website: *

Your application's publicly accessible home page, where users can go to download, make use of, or find out more information about your application. This fully-qualified URL is used in the source attribution for tweets created by your application and will be shown in user-facing authorization screens.
(If you don't have a URL yet, just put a placeholder here but remember to change it later.)

Callback URL:

Where should we return after successfully authenticating? For [@Anywhere applications](#), only the domain specified in the callback will be used. [OAuth 1.0a](#) applications should explicitly specify their `oauth_callback` URL on the request token step, regardless of the value given here. To restrict your application from using callbacks, leave this field blank.



How to Register a Twitter App in 8 Easy Steps

❖ #6 Create Your Access Token

Your access token

It looks like you haven't authorized this application for your own Twitter account yet. For your convenience, we give you the opportunity to create your OAuth access token here, so you can start signing your requests right away. The access token generated will reflect your application's current permission level.

Create my access token



How to Register a Twitter App in 8 Easy Steps

❖ #7 Choose what Access Type You Need

Application Type

Access:

- ☒ Read only
- ☐ Read and Write
- ☐ Read, Write and Access direct messages

What type of access does your application need? Note: @Anywhere applications require read & write access.

Find out more about our [Application Permission Model](#).



How to Register a Twitter App in 8 Easy Steps

❖ #8 Make a note of your OAuth Settings

OAuth settings

Your application's OAuth settings. Keep the "Consumer secret" a secret. This key should never be human-readable in your application.

Access level	Read-only About the application permission model
Consumer key	[REDACTED]
Consumer secret	[REDACTED]
Request token URL	https://api.twitter.com/oauth/request_token
Authorize URL	https://api.twitter.com/oauth/authorize
Access token URL	https://api.twitter.com/oauth/access_token
Callback URL	None
Sign in with Twitter	No

Your access token

Use the access token string as your "oauth_token" and the access token secret as your "oauth_token_secret" to sign requests with your own Twitter account. Do not share your oauth_token_secret with anyone.

Access token	[REDACTED]
Access token secret	[REDACTED]
Access level	Read-only

Recreate my access token



Exercises

❖ SPARK Streaming: StreamUtils

```
public class StreamUtils {  
  
    private static String CONSUMER_KEY = "AFiNCb8vxYZfhPls2DXyDpF";  
    private static String CONSUMER_SECRET = "JRg7SyVFkXEESWbzFzC1xaIGRC3xNdTvrekMvMFk6tjKoo0R";  
    private static String ACCESS_TOKEN = "493498548-HCt6LCposCb3Ij7Ygt7ssTxTBPwGoPrnkkDQoaN";  
    private static String ACCESS_TOKEN_SECRET = "3px3rnBzWa9bm0m0QPWNMpYc4qdOr0dxGFgp6XiCkEKH";  
  
    public static OAuthAuthorization getAuth() {  
  
        return new OAuthAuthorization(  
            new ConfigurationBuilder().setOAuthConsumerKey(CONSUMER_KEY)  
                .setOAuthConsumerSecret(CONSUMER_SECRET)  
                .setOAuthAccessToken(ACCESS_TOKEN)  
                .setOAuthAccessTokenSecret(ACCESS_TOKEN_SECRET)  
                .build());  
    }  
}
```



Exercises

❖ SPARK Streaming: StreamingOnTweets

```
public class StreamingOnTweets {  
  
    JavaStreamingContext jssc;  
  
    public JavaDStream<Status> loadData() {  
        SparkConf conf = new SparkConf()  
            .setAppName("Play with Spark Streaming");  
  
        // create a java streaming context and define the window (2 seconds batch)  
        jssc = new JavaStreamingContext(conf, Durations.seconds(2));  
  
        System.out.println("Initializing Twitter stream...");  
  
        // create a DStream (sequence of RDD). The object tweetsStream is a  
        // DStream of tweet statuses:  
        // - the Status class contains all information of a tweet  
        // See http://twitter4j.org/javadoc/twitter4j/Status.html  
        JavaDStream<Status> tweetsStream =  
            TwitterUtils.createStream(jssc, StreamUtils.getAuth());  
  
        return tweetsStream;  
    }  
}
```




Exercises

❖ SPARK Streaming: StreamingOnTweets

```
/**
 * Print the status text of the some of the tweets
 */
public void tweetPrint() {
    JavaDStream<Status> tweetsStream = loadData();

    JavaDStream<String> status =
        tweetsStream.map(tweetStatus -> tweetStatus.getText());
    status.print();

    // Start the context
    jssc.start();
    jssc.awaitTermination();
}
```



Exercises

❖ SPARK Streaming: StreamingOnTweets

```
/**
 * Find the 10 most popular Hashtag in the last minute
 */
public String top10Hashtag() {
    JavaDStream<Status> tweetsStream = loadData();

    // First, find all hashtags
    // stream is like a sequence of RDD so you can do all the operation
    // you did in the first part of the hands-on
    JavaDStream<String> hashtags = tweetsStream.
        flatMap(tweet -> Arrays.asList(tweet.getText().split(" ")))
        .filter(word -> word.matches("#(\\w+)") && word.length() > 1);
```




Exercises

❖ SPARK Streaming: StreamingOnTweets

```
// Make a "wordcount" on hashtag  
// Reduce last 60 seconds of data  
JavaPairDStream<Integer, String> hashtagMention =  
    hashtags.mapToPair(mention -> new Tuple2<>(mention, 1))  
    .reduceByKeyAndWindow((x, y) -> x + y, new Duration(60000))  
    .mapToPair(pair -> new Tuple2<>(pair._2(), pair._1()));
```



Exercises

❖ SPARK Streaming: StreamingOnTweets

```
// Then sort the hashtags  
JavaPairDStream<Integer, String> sortedHashtag =  
    hashtagMention.transformToPair(  
        hashtagRDD -> hashtagRDD.sortByKey(false));
```




Exercises

❖ SPARK Streaming: StreamingOnTweets

```
// and return the 10 most populars
List<Tuple2<Integer, String>> top10 = new ArrayList<>();

sortedHashtag.foreachRDD(rdd -> {
    List<Tuple2<Integer, String>> mostPopular = rdd.take(10);
    top10.addAll(mostPopular);

    return null;
});
```



Exercises

❖ SPARK Streaming: StreamingOnTweets

```
// we need to tell the context to start running the computation we
// have setup. It won't work if you don't add this!
jssc.start();
jssc.awaitTermination();

return "Most popular hashtag :" + top10;
}
```




Exercises

❖ SPARK Streaming: execution

```
$HOME/spark-1.6.1-bin-hadoop2.4/bin/spark-submit  
  --class "streaming.StreamingOnTweets"  
  --master local[4]  
  --packages "org.apache.spark:spark-streaming-twitter_2.10:1.5.1"  
  --jars $HOME/spark-in-practice-1.0.jar  
         $HOME/twitter4j-core-3.0.3.jar
```




DataFrame and Streaming

13/05/2019 - Big Data 2019